

THE FACT DESK LESSON • LESSON 3 OF 6

Hallucination Hunters



A 45-minute AI fact-checking lesson. The hook: students get hired at a newsroom fact desk, and nothing runs until they clear it. They learn why AI can be confidently wrong — and the four-step method that catches it.

AT A GLANCE

TIME	MATERIALS	TECH NEEDED	PREP
45 min	The Fact Desk worksheet (1 per student)	None — fully unplugged; devices only add	Print worksheets. That's it.

OBJECTIVES

- Students can explain why AI "hallucinates": it predicts likely words, so it can be confidently wrong.
- Students can identify the checkable claims in a text — names, dates, numbers, places, quotes — and rank them by risk.
- Students can lateral-read a claim (verify it against a second, independent source) and deliver an evidence-based verdict: RUN, FIX, or KILL.

STANDARDS ALIGNMENT

CCSS ELA RI.6.8

ISTE 1.3.b

ISTE 1.3.d

Full text + taught/assessed map for each code in the Teacher Guide.

🔍 WARM-UP — THE STORY THAT ALMOST RAN (7 MIN)

Read this aloud, smooth and confident, like a news anchor. Then ask for a vote: **"This is going in tomorrow's paper. Print it?"**

"Our solar system has eight planets orbiting one star, the Sun. The closest planet to the Sun is Venus, which is also the hottest. Mars is called the Red Planet because rusty iron dust covers its surface. Saturn is famous for its rings — it is the only planet in our solar system that has them. Sunlight takes about eight minutes to reach Earth."

Most classes vote print. Reveal: **two claims are false.** Mercury is closest to the Sun (Venus really is the hottest — the half-truth makes it slippery), and Saturn isn't alone: Jupiter, Uranus, and Neptune have faint rings too.

Bridge: "That paragraph sounded right because it was SMOOTH, not because it was TRUE. AI talks like that all the time — it predicts likely words, it doesn't check them. Today you're the fact desk: nothing runs until you clear it."

DIRECT INSTRUCTION — WHY AI LIES SMOOTHLY + THE METHOD (10 MIN)

Why it happens (3 min): "AI is autocomplete with confidence. It predicts the next likely word, like your phone guessing your next text. When the likely words aren't true, it says the wrong thing in the same confident voice. That's a *hallucination*." (Met the trickster genie in Lesson 1? Same trickster, new disguise.)

The method (7 min): write the four steps on the board; model steps 1–2 on the warm-up paragraph:

1. **UNDERLINE** every checkable claim: names, dates, numbers, places, quotes.
2. **RANK** — deadline reality: you can't check everything. Claims that would hurt most if wrong get **CHECK**; low-risk common knowledge gets **SAFE**.
3. **LATERAL-READ** — leave the passage; verify each CHECK claim against a second, independent source (encyclopedia, textbook, new search tab).
4. **VERDICT** — **RUN** (clean), **FIX** (correct it, then run), or **KILL** (the core claim fails — cut the story).

Say it plainly: "Don't ask the suspect if he's innocent — asking the same AI 'are you sure?' is not a check. LEAVE the answer, open a second, independent source. That's lateral reading."

ACTIVITY — THE FACT DESK (20 MIN)

Pairs become fact-check desks at *The Comet*, the school paper. Editor-in-chief Deja's memo (on the worksheet): nothing runs until the desk clears it.

1. **Desk shift 1 (7 min):** Partners work Draft #1 (Priya's Moon column) together — underline, mark C or S, log where they'd verify, stamp a verdict with evidence.
2. **Desk shift 2 (9 min):** One partner leads Draft #2 (Theo's geography quiz), the other Draft #3 (Marcus's invention feature). Swap and double-check before stamping.
3. **Verdict stand-up (4 min):** Each desk announces one verdict plus its best evidence. Tally RUN / FIX / KILL on the board; argue out disagreements.

DEBRIEF (4 MIN)

- Which planted error was hardest to catch — and what made it slippery?
- Where else does confident-but-wrong show up? Feeds, group chats, AI homework helpers...
- What's your personal "before I repeat it" rule from now on?

EXIT TICKET (4 MIN)

Bottom of worksheet: (1) in your own words, why can AI sound 100% sure and still be wrong; (2) one claim you saw online this week and the exact lateral-reading move you'd make on it.

DIFFERENTIATION

- **Support:** Assign Drafts #1–2 only; share the hint "each draft hides exactly two problems."
- **Extension:** Write your own smooth four-sentence passage with ONE planted error; trade desks and hunt.
- **ELL:** Claim types are concrete categories — students label the TYPE (name, date, number, place, quote) above each underline; partners read passages aloud.
- **With devices:** Desks lateral-read their top CHECK claims live. Without: they log where they'd verify — the method is the lesson.