

## THE FACT DESK LESSON • LESSON 3 OF 6

# Hallucination Hunters



A 45-minute AI fact-checking lesson. Students run the fact desk at the school paper, and nothing gets printed until they clear it. They learn why AI can sound sure and still be wrong, plus a four-step method for catching it. Stands alone.

## AT A GLANCE

| TIME   | MATERIALS                                                                                                                  | TECH NEEDED                                                                    | PREP                         |
|--------|----------------------------------------------------------------------------------------------------------------------------|--------------------------------------------------------------------------------|------------------------------|
| 45 min | The Fact Desk worksheet (3 pages, 1 per student) · Slides (optional) · the board, for the RUN / FIX / KILL tally · pencils | None. With devices, desks open the source they named and check their C claims. | Print worksheets. That's it. |

## OBJECTIVES — STUDENTS CAN...

- Explain why AI "hallucinates": it predicts likely words, so it can be confidently wrong. (*Exit #1*)
- Identify the checkable claims in a text (names, dates, numbers, places, quotes) and rank them by risk. (*drafts 1–3, C/S*)
- Lateral-read a claim (name a second, independent source that could confirm it) and give an evidence-based verdict: RUN, FIX, or KILL. (*second-source column + stamp*)

## STANDARDS ALIGNMENT

CCSS ELA RI.6.8

ISTE 1.3.b

ISTE 1.3.d

Full text + taught/assessed map for each code in the Teacher Guide.

## DIFFERENTIATION

- **Support:** Drafts #1–2 only. In Round 2, both partners take Draft #2 and swap as normal. Hint: "Each draft hides two false facts, plus a shaky ending."
- **Extension:** Write a smooth four-sentence passage with ONE planted error, or none at all. Trade. Your partner has to decide: RUN, FIX, or KILL.
- **ELL:** Label the claim TYPE (name, date, number, place, quote) above each underline; partners read aloud; desk words are glossed on p.1.
- **No-device classrooms:** built in. Students name the source they would check. **With devices,** they open it.

## WARM-UP — THE STORY THAT ALMOST RAN (6 MIN)

Read this aloud, smooth and confident, like a news anchor. Then call the vote: **"This is going in tomorrow's paper. Print it?"**

*"Our solar system has eight planets orbiting one star, the Sun. The closest planet to the Sun is Venus, which is also the hottest. Mars is called the Red Planet because rusty iron dust covers its surface. Saturn is famous for its rings — it is the only planet in our solar system that has them. Sunlight takes about eight minutes to reach Earth."*

Expect most of the room to vote print. Reveal: **two sentences are false.** Mercury is closest to the Sun (Venus really is the hottest — the half-truth is what makes it slippery), and Saturn isn't alone: Jupiter, Uranus, and Neptune have faint rings too.

*If the room votes HOLD:* ask which sentence they doubt — most name Saturn and miss Venus. "One doubt is a start. The desk finds BOTH." *If a kid names Venus instantly:* "Write it down. We vote first, then we check." The second error usually slips past anyway.

**Bridge (say it):** "That sounded right because it was SMOOTH, not because it was TRUE. AI predicts likely words, and it doesn't always check them. So it can sound exactly like that. Today you're the fact desk: nothing gets printed until you clear it."

## DIRECT INSTRUCTION — WHY AI SOUNDS SURE + THE METHOD (10 MIN)

**Why it happens (3 min):** Play class autocomplete. "Finish my sentence. Twinkle, twinkle, little..." (Star.) "Bulls get mad when they see the color..." (Red.) "You just did what AI does: you picked the likely word. The first one was right. The second wasn't. Bulls can't see red well; it's the moving cape. AI is autocomplete with confidence. When the likely words aren't true, it says the wrong thing in the same confident voice. That's a *hallucination*." (Did Lesson 1? "The genie gives you what you SAID. Autocomplete gives you what's LIKELY. Neither one promises what's TRUE.")

*If someone asks "so is AI always wrong?": "No, it's right a lot. That's the danger: you can't tell by the tone." If "my AI shows links": "Good. Click one. A link counts as a second source only if you open it and it says the same thing."*

**The method (7 min):** Steps up (slide 5). Then model steps 1 and 2 out loud on the corrected warm-up (slide 6).

1. **UNDERLINE** every checkable claim: names, dates, numbers, places, quotes.
2. **RANK** — you can't check everything. Would it hurt if wrong? → **C** (check). Low risk (not a record, not a number, and the ending doesn't need it) → **S** (safe). "Everybody knows it" is not a reason.
3. **LATERAL-READ** — leave the passage; name a second, independent source that could confirm each C.
4. **VERDICT** — **RUN** (clean) · **FIX** (correct it, then print) · **KILL** (core claim fails: cut it). Does the ending stand?

**Say it plainly (slide 7):** "Don't ask the suspect if they're innocent. Asking the same AI 'are you sure?' is not a check. LEAVE the answer and open a second, independent source. That's lateral reading."

*If a student says "I'd just Google it": "Google shows you ten sources. Which one, and why that one? That's your log."*

**Hand-out (say it):** "Pair up. Worksheets out. Deja's note is at the top of page 1. Draft #1 together, seven minutes."

## ACTIVITY — THE FACT DESK (21 MIN)

Pairs are fact desks at *The Comet*, the school paper. Play the editor, not the answer key.

1. **Round 1 (7 min):** Both partners work Draft #1 (Priya's Moon column) together, each on their OWN sheet: underline, mark C or S, log their two riskiest claims, pick a stamp, write the note with evidence.  
*If a desk asks "is this one true?": "That's a checkable claim. Where's your second source?"*  
*Transition:* "Round change. Partner A is the partner whose first name comes first in the alphabet. A, take Theo's draft on page 2. B, take Marcus's on page 3. Six minutes, your own sheet. Stamp and note before the swap."
2. **Round 2 (10 min):** Solo, all four steps, stamp and note included. (6 min) Then: **"Swap sheets. Read your partner's log, add one claim they missed in an empty row, and hand it back."** (2 min) Then: "Take your sheet back. If the new claim changes your stamp or note, fix it." (2 min)  
*If a desk finishes early:* the Extension on page 1. *If an odd number of students:* one desk of three, and the third takes Draft #3 too. At the swap, each person passes their sheet to the left, then sheets come straight back.  
*Transition:* "Pens down. Hands up for your stamps first, then I call on desks for evidence."
3. **Verdicts out loud (4 min):** Draft by draft, hands up for RUN, FIX, and KILL. You tally one row per draft on the board. Then, for each stamp that got hands, one desk gives its stamp and best evidence: five or six desks in all.  
*If two desks disagree:* "Score the evidence, not the stamp. Who named a claim, a fix, AND a second source?"  
*Transition:* "Stamps are in. Now the part that matters after today: where else does this show up?"

## DEBRIEF (4 MIN)

- Which planted error was hardest to catch, and what made it slippery?
- Where else does confident-but-wrong show up? Feeds, group chats, AI homework helpers...
- What's your personal "before I repeat it" rule from now on?

*If nobody names the fake quote:* "Did anyone check the Edison quote? Why not?" Quote marks feel like proof.

**Transition (say it):** "Bottom of page 3: Last Check. Solo, four minutes. Two questions, then hand it in."

## EXIT TICKET — LAST CHECK (4 MIN)

Worksheet p.3, solo: the two Last Check questions, word for word as printed. Collect; keys in the Teacher Guide. Each sheet shows Draft #1 plus #2 or #3. That's by design.

*If a student writes "Google it" for #2:* same push-back as in Direct Instruction. "Which one, and why that one?"